



Just Relax It

Discrete variables relaxation

Daniil Dorin, Igor Ignashin, Nikita Kiselev, Andrey Veprikov

Intelligent Systems, MIPT, 2024

Project description

Motivation


For lots of mathematical problems we need an ability to sample discrete random variables. But the usage of truly discrete random variables is infeasible. Thus we use different relaxation methods.

Used algorithms


- | | |
|---------------------------------|------------------------|
| 1. Relaxed Bernoulli | 5. Invertible Gaussian |
| 2. Correlated relaxed Bernoulli | 6. Hard concrete |
| 3. Gumbel-softmax TOP-K | 7. REINFORCE |
| 4. Straight-Through Bernoulli | 8. Logit-Normal |

Solution

Just Relax It is a flexible, scalable deep probabilistic programming library built on PyTorch and Pyro.



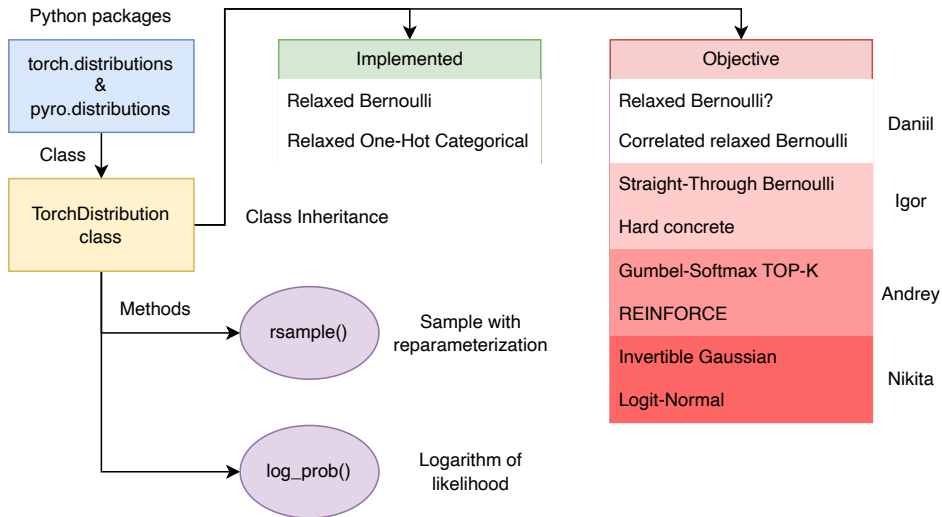
Discrete



Gumbel-Softmax 1/7

Scheme of the project

Project scheme



Brief algorithms description

Straight-Through Bernoulli, Hard concrete (Igor)

Straight-Through Bernoulli

$$u \sim \mathcal{U}[0,1]$$

$$s = \sigma\left(\frac{\log u + \log(1-u) + \log \alpha}{\beta}\right)$$

$$\bar{s} = s(\zeta - \gamma) + \gamma$$

$$z = \min(1, \max(0, \bar{s}))$$

Hard concrete

$$p_i = \sigma(a_i)$$

$$b_i \sim \text{Binomial}(\sqrt{p_i})$$

$$h_i = b_i \sqrt{p_i}$$

$$\mathbb{E} h_i = p_i$$

$$\mathbb{E} f(h_i) = f(p_i) + o(\sqrt{p_i})$$

Invertible Gaussian, Logit-Normal (Nikita)

Invertible Gaussian Reparameterization

Remove interpretability in Gumbel-Softmax

Objective: relax one-hot $\mathbf{z} \sim \text{Cat}(\boldsymbol{\pi})$

GS: $\tau \rightarrow 0$ concentrates mass on vertices:

$$\tilde{\mathbf{z}} = \text{softmax}\left(\frac{\log \boldsymbol{\pi} + \mathbf{G}}{\tau}\right), \mathbf{G}_i \sim \text{Gumbel}(0, 1)$$


IGR: map $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ to simplex, using invertible $g(\cdot, \tau)$ with temperature τ :

$$\mathbf{y} = \boldsymbol{\mu} + \text{diag}(\boldsymbol{\sigma})\boldsymbol{\epsilon},$$


$$\tilde{\mathbf{z}} = g(\mathbf{y}, \tau) = \text{softmax}_{++}(\mathbf{y}/\tau)$$

Logit-Normal distribution

$\text{softmax}(\mathbf{x})$ of $\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$



Dirichlet




Logistic Normal

- $\text{Logit-Normal}(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \rightarrow \text{Dir}(\boldsymbol{\alpha})$
- $\text{Dir}(\boldsymbol{\alpha}) \rightarrow \text{Logit-Normal}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$
- KL divergence?

Gumbel-softmax TOP-K, REINFORCE (Andrey)

Gumbel-softmax TOP-K




GS: one sample from $\text{Cat}(\alpha)$

GS TOP-K: sample K elements without replacement via

$$i_1, \dots, i_K = \arg \text{topK}_i \{G_i + \log \alpha_i\}$$

REINFORCE

- Score function (SF) vs pathwise derivative (PD) estimator of $\frac{\partial}{\partial \theta} \mathbb{E}_x[f(x)]$



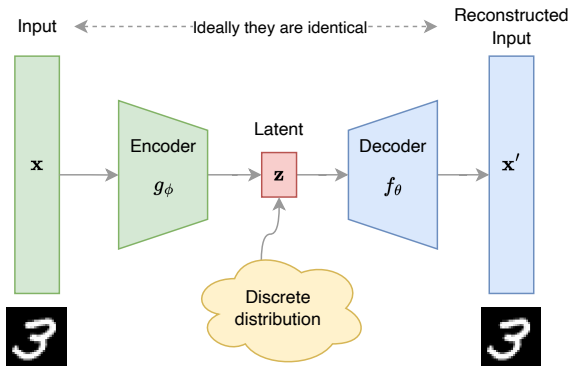
- Monte-Carlo policy gradient estimation of $\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi} [G_t \nabla_{\theta} \ln \pi_{\theta}(A_t | S_t)]$:

- $\pi_{\theta} \rightarrow S_1, A_1, R_2, S_2, A_2, \dots, S_t$
- $\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \ln \pi_{\theta}(A_t | S_t)$

Idea for demo/basic code

Demo/basic code

- Consider a VAE architecture
- Classical latent space is Gaussian
- Discrete latent space needs gradient computation
- Implement^a different relaxation methods for these discrete latents



^aSee our first attempts [here](#)